



# Solution - 4


## Revision 2


1. The solutions might be different when using different split rules.

All reasonable solution will get points.



2.





3.

For maximal height of a 2-4 tree, we will be having one key per node, hence it will behave like a Binary Search Tree.

keys at level 0 = 1

keys at level 1 = 2

keys at level 2=4 and so on. . .

Adding total number of keys at each level we get a GP on solving which we will get the maximal height of the tree.

Hence, height =  $\log_2(n + 1) - 1$

For minimal height of a 2-4 tree, we will be having three keys(maximum possible number) per node. keys at level 0 = 3

keys at level 1 = 3\*(4)

keys at level 2=3\*(42)and so on. . .

Hence, height =  $\log_4(n + 1) - 1$



4.

The shortest simple path from any node x will be the black height of the tree with x as root(i.e.,  $bh(x)$ ). There could be many branches in the tree; each branch is a combination of red and black nodes. The longest simple path in any tree will be that path which has the total number of nodes = (Property 4)  $bh(x) + \text{max possible number of red nodes}$ . The maximum possible number of red

nodes will be equal to the  $bh(x)$ , as to satisfy the red-black property., for each red node, its children has to be black (no two consecutive red nodes in a path). Hence the max height of the tree could be  $2 * bh(x)$ , twice the shortest simple path.


5.


Case 1: Sibling of parent is red



Case 2 and 3:

These two cases should be considered together, because case 2 leads to case 3.





6.

In this problem, we want to use the first fit heuristic to find the smallest drive that would fit.

We can take our USB drives and convert them into red-black tree. Since there are  $n$  drives, we can build this tree in  $O(n)$  time. We then loop through our images and for each image we perform the following steps:

1. We find the first drive that has enough space to hold the image, which takes  $O(\log n)$  time.
2. We delete this node from our tree, which takes  $O(\log n)$  time.
3. We then insert this node back into our tree in  $O(\log n)$  time.

For each image we spend  $O(\log n + \log n + \log n) = O(\log n)$  time. For processing all  $m$  images, the overall time complexity is  $O(n + m \log n) = O(m \log n)$ .

Pseudocode:

StoreImage(drives, images):

tree = BUILD\_RED\_BLACK\_TREE(drives) //  $O(n)$

for each image in images:

drive = SEARCH\_FIRST\_FIT(tree, image.size) //  $\log(n)$

DELETE\_FROM\_TREE(tree, drive) //  $\log(n)$

drive.size = drive.size - image.size

```
        ADD_TO_TREE(tree, drive)                // log(n)
    }
```

7.

based on the properties of BTS, we can pick the middle element( $\text{nums}[(0+\text{nums.length})/2]$ ) from the sorted array and it is the root of tree. Based on this, we can use recursion to do the same thing on elements from 0 to mid, and (mid+1) to  $\text{nums.length}$

```
public TreeNode sortedArrayToBST(int[] nums) {

    return recursion(nums,0,nums.length);

}

public TreeNode recursion(int[] nums,int i,int j){

    if(i>=j) return null;

    TreeNode root=new TreeNode(nums[(i+j)/2]);

    root.left=recursion(nums,i,(i+j)/2);

    root.right=recursion(nums,(i+j)/2+1,j);

    return root;

}
```

8.

11. (a) Assume a binary tree has height of  $h$ , and for each node  $\max(\text{size}(t.\text{left}), \text{size}(t.\text{right})) \leq (\frac{2}{3})\text{size}(t)$ . Then we have  $(\frac{2}{3})^h n \geq 1$ . Hence,  $h \leq \log_{\frac{2}{3}} \frac{1}{n} = O(\log n)$ . This satisfies the definition of a balanced binary tree.

(b) 1: **procedure** LEFT-ROTATE( $T, x$ )  
 2:    $y = x.\text{right}$   
 3:    $\text{temp} = x.\text{size}$   
 4:   **if**  $y.\text{left} = \text{NIL}$  **then**  
 5:      $x.\text{size} = x.\text{size} - y.\text{size}$   
 6:   **else**  
 7:      $x.\text{size} = x.\text{size} - y.\text{size} + y.\text{left}.\text{size}$   
 8:    $y.\text{size} = \text{temp}$   
 9:   LEFT-ROTATE( $T, x$ ) // From lecture  
 1: **procedure** RIGHT-ROTATE( $T, x$ )  
 2:    $y = x.\text{left}$   
 3:    $\text{temp} = x.\text{size}$   
 4:   **if**  $y.\text{right} = \text{NIL}$  **then**  
 5:      $x.\text{size} = x.\text{size} - y.\text{size}$   
 6:   **else**  
 7:      $x.\text{size} = x.\text{size} - y.\text{size} + y.\text{right}.\text{size}$   
 8:    $y.\text{size} = \text{temp}$   
 9:   RIGHT-ROTATE( $T, x$ ) // From lecture

- (c) Find the median node and apply rotation so that this node is the root of the tree. Then apply the same process recursively to its children.

1: **procedure** PERFECT-BALANCE( $T$ )  
 2:   Balance-Tree( $T.\text{root}$ )  
 1: **procedure** BALANCE-TREE( $T$ )  
 2:    $n = T.\text{root}$   
 3:    $\text{medianValue} = n.\text{size}/2$   
 4:    $\text{medianNode} = \text{FindMedian}(n, \text{medianValue})$   
 5:    $\text{parent} = n.p$   
 6:   **while**  $T.\text{root} \neq \text{medianNode}$  **do**  
 7:     **if**  $T.\text{left} \neq \text{null}$  **and**  $T.\text{left}.\text{size} \geq \text{medianValue}$  **then**

```

8:         right-rotate(T,medianNode.p)
9:     else if T.right  $\neq$  null and T.right.size  $\geq$  medianValue then
10:         left-rotate(T,medianNode.p)
11: if parent  $\neq$  null then
12:     if parent.left == n then
13:         parent.left = medianNode
14:     else
15:         parent.right = medianNode
16: medianNode.p = parent
17: if medianNode.left  $\neq$  null then
18:     Balance-Tree(medianNode.left)
19: if medianNode.right  $\neq$  null then
20:     Balance-Tree(medianNode.right)
21: return

1: procedure FIND-MEDIAN(NODE ROOT, INT VALUE)
2:     if root.left.size + 1 == value then
3:         return root
4:     else if root.left.size geq value then
5:         return FindMedian(n.left,value)
6:     else
7:         return FindMedian(root.right, value-1-root.left-size)

```

9. The following solution has been adapted from an anonymous student's solution.

Let the system be implemented using a two level data-structure of nested red black trees. The first level is a single red-black tree where the keys of the nodes are all the possible first characters of the names that are to be stored into the tree and each node, contains a pointer, say *lvl2*, to another red-black tree which is used to actually store all names (and other data) with identical starting letters.

Let the structure of a level 1 node and a level 2 node be as follows:

class level_1_node	class level_2_node
1    Initialization( <i>character</i> )	1    Initialization( <i>name, data</i> )
2 <i>key</i> = <i>character</i>	2 <i>key</i> = <i>name</i>
3 <i>level2</i> = NULL	3 <i>left</i> = NULL
4 <i>left</i> = NULL	4 <i>right</i> = NULL
5 <i>right</i> = NULL	5 <i>p</i> = NULL
6 <i>p</i> = NULL	6 <i>Color</i> = RED
7 <i>color</i> = RED	7 <i>data</i> = <i>data</i>

Let *T* be the root node of the whole structure and is of a root of the level 1 red black tree.

- 1) Insert: when we are to insert a new name into the system, we first check if the first character of the new name already exists in the tree rooted at *T*.
  - a) If it does exist, we get a pointer to that node, say *level1\_ptr* by RB-SEARCH. Using *level1\_ptr* we get the root of the level2 red black tree (which contains the all names with the same first character the name we wish to insert) and use RB-INSERT to insert the new name into the level2 red black tree.
  - b) If it does not exist, we first create a new red black tree with root, say *level2\_root* and use RB-INSERT to insert the new name into the tree. Subsequently we create a new level1 tree node with the first character as the key and set the pointer *lvl2* to *level2\_root*. We then use RB-INSERT to insert the level1 node into the first level red-black tree

Time complexity  $O(\log n)$



Insert Pseudo code:

```
Insert(T, name, data)
1  lvl2node = new level_2_node(name, data)
2
3  first_character = name[1]
4  level1_ptr = RB-SEARCH(T, first_character)
5
6  if level1_ptr == NULL
7      lvl1node = new level_1_node(first_character)
8
9      level1_ptr = lvl1node
10     RB-INSERT(T, lvl1node)
11
12     if level1_ptr.level2 == NULL
13         level1_ptr.level2 = lvl2node
14     else
15         RB-INSERT(level1_ptr.level2, lvl2node)
```

- 2) Search: we first search for the node in the level 1 tree, *lvl1node*, corresponding to the first character of the name we are looking for, i.e. RB-SEARCH(*T*, *first\_character*). Once we find the *lvl1node*, we search for the the name and get a pointer to the node containing the name and its corresponding data. RB-SEARCH(*lvl1node.level2*, *name*). Time complexity:  $O(\log n)$

Search Pseudo code:

```
Search(T, name)
1  first_character = name[1]
2  lvl1node = RB-SEARCH(T, first_character)
3
4  return RB-SEARCH(lvl1node.level2, name)
```

- 3) Delete: we first search for the node in the level1 tree, *lvl1node*, corresponding to the first character of the name we want to delete, i.e. *RB-SEARCH(T, first\_character)*. Once we find the *lvl1node*, we Delete the name from the tree rooted at *lvl1node.level2* i.e. *RB-DELETE(lvl1node.level2, name)*. Time complexity:  $O(\log n)$

Delete Pseudo code:

Delete(*T, name*)

- 1 *first\_character* = *name*[1]
- 2 *lvl1node* = *RB-SEARCH(T, first\_character)*
- 3 *RB-DELETE(lvl1node.level2, name)*

- 4) Print names starting character: we first search for the node in the level1 tree, *lvl1node*, where the key is the equal to the given character, i.e. *RB-SEARCH(T, first\_character)*. Once we have found the *lvl1node*, we perform an inorder traversal on the tree rooted at *lvl1node.level2*. Time complexity:  $O(g + \log n)$

Character Pseudo code:

GET-FIRST-NAME-SEARCH (*T, character*)

- 1 *first\_character* = *name*[1]
- 2 *lvl1node* = *RB-SEARCH(T, first\_character)*  
*name\_list* = init list
- 3 *INORDER-TRAVERSAL (lvl1node.level2, name\_list )*
- 4 **return** *name\_list*

INORDER-TRAVERSAL (*node, name\_list*)

- 1 **if** *node* == NULL
- 2 **return** NULL
- 3 *INORDER-TRAVERSAL (node.left, name\_list)*
- 4 *name\_list.insert\_head(node.key)*
- 5 *INORDER-TRAVERSAL (node.right, name\_list)*

10. To solve this question, let us use a Red-Black tree where each node is augmented with 2 pieces of extra information.

- 1) every node stores the size of its subtree in an attribute called **size** with default of 1
- 2) every node stores the sum of all rates of its subtree, i.e. the sum of the rates of all nodes present in the subtree including the node itself in an attribute called **rate\_sum** default of the node's own rate

**a)** Insert(*d*, *r*): This is very similar to a regular RB tree insert which takes  $O(\log n)$ , where the key is the distance *d*. We create a new node, with attributes distance, rate, size and rate\_sum; where distance = *d*, rate = *r*, size = 1 and rate\_sum = *r*. Slight modifications need to be made to RB-INSERT to maintain our augmentations. As we search for the correct position to insert our new node

- 1) we increment the **size** attribute of every node we encounter along the path.
- 2) we add the new node's rate attribute to the **rate\_sum** attribute of every node we encounter along the path

RB-INSERT(*T*, *z*)

```

1  y = T.nil
2  x = T.root
3  while x ≠ T.nil
4      y = x
5      x.size = x.size + 1
6      x.rate_sum = x.rate_sum + z.rate
7      if z.key < x.key
8          x = x.left
9      else
10         x = x.right
11  z.p = y
12  if y == T.nil
13      T.root = z
14  elseif z.key < y.key
15      y.left = z
16  else
17      y.right = z
18  z.left = T.nil
19  z.right = T.nil
20  z.color = RED
21  RB-INSERT-FIXUP(T, z)

```

class RB-NODE

```

1  Initialization(d, r)
2      key = d
3      distance = d
4      rate = r
5      size = 1
6      rate_sum = rate
7
8      p = null
9      left = null
10     right = null
11     color = null

```

INSERT(*d*, *r*) //assuming *T* is accessible

```

1  node = new RB-NODE(d, r)
2  RB-INSERT(T, node)

```

We also need to modify the LEFT-ROTATE and RIGHT-ROTATE operations to maintain the augmentations

LEFT-ROTATE (T, x)

```

1  y = x.right
2  x.right = y.left
3  if y.left ≠ T.nil
4      y.left.p = x
5  y.p = x.p
6  if x.p == T.nil
7      T.root = y
8  else if x == x.p.left
9      x.p.left = y
10 else
11     x.p.right = y
12 y.left = x
13 x.p = y
14 y.size = x.size
15 y.rate_sum = x.rate_sum
16 x.size = x.left.size + x.right.size + 1
17 x.rate_sum = x.left.rate_sum +
                x.right.rate_sum +
                x.rate

```

RIGHT-ROTATE (T, x)

```

1  y = x.left
2  x.left = y.right
3  if y.right ≠ T.nil
4      y.right.p = x
5  y.p = x.p
6  if x.p == T.nil
7      T.root = y
8  else if x == x.p.right
9      x.p.right = y
10 else
11     x.p.left = y
12 y.right = x
13 x.p = y
14 y.size = x.size
15 y.rate_sum = x.rate_sum
16 x.size = x.left.size + x.right.size + 1
17 x.rate_sum = x.left.rate_sum +
                x.right.rate_sum +
                x.rate

```

**b)** delete(k): we first, need to find the node that is at  $k^{\text{th}}$  distance away from the headquarters. We can do this in  $O(\log n)$  using OS-SELECT(T, k) where T is the root. We can use OS-SELECT because the nodes of our tree contain the size augmentation that is needed by OS-SELECT. Let the  $k^{\text{th}}$  node be knode. After finding knode, we can find it's key and toll rate. We then search for the knode using it's key (similar to searching for a node in a binary tree) and decrement the size by 1 and the rate\_sum by knode.rate for all nodes we encounter along the path as we search for the knode in  $O(\log n)$ . Then we invoke the subroutine RB-DELETE(T, knode) where T is the root node and knode is the node that we want to delete which takes  $O(\log n)$ .

**c)** The core logic behind this method is that; since every node maintains the size of it's subtree and also the sum of all rates in the subtree, we find the root of the subtree of interest (this subtree must contain both the lower bound and the upper bound nodes) and remove the branches of this subtree that we don't require (branches that contain keys not in the lower and upper bounds), we would obtain the correct number of number of nodes and the sum of rates of all nodes within the lower and upper bounds. Thus the average is the sum of all rates / number of nodes.

The method FIND-LEFT-INCORRECT and FIND-RIGHT-INCORRECT are used to find the sum of rates and sum of sizes of those nodes that are not within the Lower bound and Upper bound respectively.

The method GET-LOWEST-COMMON-ANCESTOR is used to get the root of the subtree that contains both the lower bound node and the upper bound node.

FIND-LEFT-INCORRECT (*node*, *Lbound*, *minus\_rate\_sum*, *minus\_size\_sum*)

```
1  if node.key == Lbound
2      if node.left != NULL
3          minus_rate_sum = minus_rate_sum + node.left.rate_sum
4          minus_size_sum = minus_size_sum + node.left.size
5
6  else if node.key < Lbound
7      minus_rate_sum = minus_rate_sum + node.rate
8      minus_size_sum = minus_size_sum + 1
9
10     if node.left != NULL
11         minus_rate_sum = minus_rate_sum + node.left.rate_sum
12         minus_size_sum = minus_size_sum + node.left.size
13
14     FIND-LEFT-INCORRECT(node.right, Lbound, minus_rate_sum, minus_size_sum)
15
16 else
17     FIND-LEFT-INCORRECT(node.left, Lbound, minus_rate_sum, minus_size_sum)
```

FIND-RIGHT-INCORRECT (*node*, *Ubound*, *minus\_rate\_sum*, *minus\_size\_sum*)

```
1  if node.key == Ubound
2      if node.right != NULL
3          minus_rate_sum = minus_rate_sum + node.right.rate_sum
4          minus_size_sum = minus_size_sum + node.right.size
5
6  else if node.key > Ubound
7      minus_rate_sum = minus_rate_sum + node.rate
8      minus_size_sum = minus_size_sum + 1
9
10     if node.right != NULL
11         minus_rate_sum = minus_rate_sum + node.right.rate_sum
12         minus_size_sum = minus_size_sum + node.right.size
13
14     FIND-RIGHT-INCORRECT(node.left, Ubound, minus_rate_sum, minus_size_sum)
15
16 else
17     FIND-RIGHT-INCORRECT(node.right, Ubound, minus_rate_sum, minus_size_sum)
```

GET-LOWEST-COMMON-ANCESTOR (*T*, *Lbound*, *Ubound*)

```
1  pointer_node = T.root
2  while True
3      if pointer_node.key < Lbound and pointer_node.key < Ubound
4          pointer_node = pointer_node.right
5      else if pointer_node.key > Lbound and pointer_node.key > Ubound
6          pointer_node = pointer_node.left
7      else
8          break
9  return pointer_node
```

TOLL(*i*, *j*) //assuming the RB tree 'T' is accessible to this method

```
1  inode = OS-SELECT(T, i) // finding the ith node. //O(log n)
2  jnode = OS-SELECT(T, j) // finding the jth node. //O(log n)
3  idist = inode.distance
4  jdist = jnode.distance
5  lca = GET-LOWEST-COMMON-ANCESTOR(T, idist, jdist) //O(log n)
6
7  minus_rate_sum = 0
8  minus_size_sum = 0
9
10 FIND-LEFT-INCORRECT(lca, idist, minus_rate_sum, minus_size_sum) //O(log n)
11 FIND-RIGHT-INCORRECT(lca, jdist, minus_rate_sum, minus_size_sum) //O(log n)
12
13 total_rate_sum = lca.rate_sum – minus_rate_sum
14 total_size_sum = lca.size – minus_size_sum
15
16 return total_rate_sum / total_size_sum
```