Introduction to Machine Learning
Problems Unit 3: Multiple Linear Regression

Prof. Sundeep Rangan

1. There is no one single correct answer to this problem. Below are possible ideas.

   (a) A possible target variable, y, is the total sales (in some units like dollars or units) for the product.

   (b) One way the features can be represented is to let x_1 be the numeric score and $x_2, x_3, \ldots, x_k$ be the frequency of occurrence for each of the $k-1$ words. The linear model could then be
   $$y = \beta_1 x_1 + \ldots + \beta_k x_k + \epsilon.$$

   (c) One way is to normalize the scores so that $x_1 = \text{score}/5$ when the score is out of five and $x_1 = \text{score}/10$ when it is out of ten.

   (d) For this case, you could use a variation of one-hot coding. Specifically, we replace the single variable x_1 with four variables $x_1, \ldots, x_4$ using the following encoding:

Review type	x_1	x_2	x_3	x_4
Score available with score $s = 1, \ldots, 5$	s	0	0	0
Good rating	0	1	0	0
Bad rating	0	0	1	0
No rating	0	0	0	1

   Then, each rating method obtains a different offset where the offset for the case of the score is proportional to the score.

   (e) Probably, fraction of reviews with "good" is a more likely to be a useful statistic than total number of reviews. For example, a very bad product could have a 1000 reviews with only 10 reviews saying "good", while a second product could have 8 good reviews out of 9. If we use total number of reviews with "good", the first product would have a higher score in that feature.

2. (a) Since there are two features,
   $$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \epsilon.$$

   (b) The transformed feature matrix and response vector is
   $$\mathbf{A} = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \\ 1 & 1 & 1 \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} 1 \\ 4 \\ 3 \\ 7 \end{bmatrix},$$

The solution is $\hat{\boldsymbol{\beta}} = (\mathbf{A}^\mathsf{T}\mathbf{A})^{-1}\mathbf{A}^\mathsf{T}\mathbf{y}$. This can be solved in python with

```
A = np.array([[1,0,0],[1,0,1],[1,1,0],[1,1,1]])
y = np.array([1,4,3,7])
beta = np.linalg.lstsq(A,y)[0]
```

which returns the solution $\hat{\boldsymbol{\beta}} = [0.75, 2.5, 3.5]$.

3. (a) Model $\hat{y} = (a_1 x_1 + a_2 x_2)e^{-x_1-x_2}$. Use the parameters and basis functions,

$$\beta = [a_1, a_2], \quad \phi(\mathbf{x}) = [x_1 e^{-x_1-x_2}, x_2 e^{-x_1-x_2}].$$

(b) The mode is

$$\hat{y} = \begin{cases} a_1 + a_2 x & \text{if } x < 1 \\ a_3 + a_4 x & \text{if } x \geq 1. \end{cases}$$

We can write this as,

$$\hat{y} = (a_1 + a_2 x)\mathbb{1}_{\{x<1\}} + (a_3 + a_4 x)\mathbb{1}_{\{x\geq 1\}},$$

where $\mathbb{1}_{\{x<1\}}$ is the *indicator function*,

$$\mathbb{1}_{\{x<1\}} = \begin{cases} 1 & \text{if } x < 1 \\ 0 & \text{if } x \geq 1. \end{cases}$$

The indicator function is very useful for models with cases like this. Then, we take

$$\beta := [a_1, a_2, a_3, a_4]$$
$$\phi(x) := \left[\mathbb{1}_{\{x<1\}}, x\mathbb{1}_{\{x<1\}}, \mathbb{1}_{\{x\geq 1\}}, x\mathbb{1}_{\{x\geq 1\}}\right].$$

(c) $\hat{y} = (1 + a_1 x_1)e^{-x_2+a_2}$. Rewrite this as,

$$\hat{y} = e^{-a_2}(1 + a_1 x_1)e^{-x_2}.$$

So, we can use the parameters and basis functions,

$$\beta = [e^{-a_2}, a_1 e^{-a_2}], \quad \phi(\mathbf{x}) = [e^{-x_2}, x_1 e^{-x_2}].$$

To get the parameters $\boldsymbol{\beta}$ back from the vector $\mathbf{a}$, we invert the equations:

$$\beta_1 = e^{-a_2} \Rightarrow a_2 = -\ln(\beta_1).\beta_2 \qquad = a_1 e^{-a_2} = a_1\beta_1 \Rightarrow a_1 = \beta_2/\beta_1.$$

4. An automobile engineer wants to model the relation between the accelerator control and the velocity of the car. The relation may not be simple since there is a lag in depressing the accelerator and the car actually accelerating. To determine the relation, the engineers measures the acceleration control input x_k and velocity of the car y_k at time instants $k = 0, 1, \ldots, T-1$. The measurements are made at some sampling rate, say once every 10 ms. The engineer then wants to fit a model of the form

$$y_k = \sum_{j=1}^{M} a_j y_{k-j} + \sum_{j=0}^{N} b_j x_{k-j} + \epsilon_k, \tag{1}$$

for coefficients a_j and b_j. In engineering this relation is called a *linear filter* and it statistics it is called an *auto-regressive moving average (ARMA)* model.

(a) The parameter vector would be

$$\boldsymbol{\beta} = [a_1, \ldots, a_M, b_0, \ldots, b_N]^\mathsf{T}.$$

Note the transpose to make this a column vector. There are $M + N + 1$ unknown parameters.

(b) We can rewrite the data as

$$\mathbf{A} = \begin{bmatrix} y_{M-1} & \cdots & y_0 & x_M & \cdots & x_{M-N} \\ y_M & \cdots & y_1 & x_{M+1} & \cdots & x_{M-N+1} \\ \vdots & \cdots & \vdots & \vdots & \cdots & \vdots \\ y_{T-2} & \cdots & y_{T-M-1} & x_{T-1} & \cdots & x_{T-N-1} \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y_M \\ \vdots \\ y_{T-1} \end{bmatrix}.$$

(c) First, we partition the matrix $\mathbf{A}$ into two parts:

$$\mathbf{A} = [\mathbf{A}_Y \ \mathbf{A}_X], \quad \mathbf{A}_Y = \begin{bmatrix} y_{M-1} & \cdots & y_0 \\ y_M & \cdots & y_1 \\ \vdots & \cdots & \vdots \\ y_{T-2} & \cdots & y_{T-M-1} \end{bmatrix}, \quad \mathbf{A}_X = \begin{bmatrix} x_M & \cdots & x_{M-N} \\ x_{M+1} & \cdots & x_{M-N+1} \\ \vdots & \cdots & \vdots \\ x_{T-1} & \cdots & x_{T-N-1} \end{bmatrix}.$$

Now,

$$\frac{1}{T}\mathbf{A}^\mathsf{T}\mathbf{A} = \begin{bmatrix} \mathbf{A}_Y^\mathsf{T}\mathbf{A}_Y & \mathbf{A}_Y^\mathsf{T}\mathbf{A}_X \\ \mathbf{A}_X^\mathsf{T}\mathbf{A}_Y & \mathbf{A}_X^\mathsf{T}\mathbf{A}_X \end{bmatrix}.$$

The first term can be simplified as

$$\frac{1}{T}(\mathbf{A}_Y^\mathsf{T}\mathbf{A}_Y)_{i,j} \overset{(a)}{=} \frac{1}{T}\sum_k (\mathbf{A}_Y^\mathsf{T})_{ik}(\mathbf{A}_Y)_{kj}$$

$$\overset{(b)}{=} \frac{1}{T}\sum_k (\mathbf{A}_Y)_{ki}(\mathbf{A}_Y)_{kj}$$

$$\overset{(c)}{=} \frac{1}{T}\sum_k y_{M+k-i}y_{M+k-j} \overset{(d)}{=} \frac{1}{T}\sum_k y_k y_{k+i-j} \approx R_{yy}(i-j),$$

where (a) follows from the definition of matrix multiplication; (b) uses the definition of a transpose matrix; (c) is the formula for the entries in $\mathbf{A}_Y$ (d) follows from a change of variables and (e) is the definition of the auto-correlation. The approximation is that there is a slight change in the end points of the summation which will not matter for large T. Similarly, one can show that

$$\frac{1}{T}(\mathbf{A}_Y^\mathsf{T}\mathbf{A}_X)_{i,j} = \frac{1}{T}\sum_k\sum_k y_{M+k-i}x_{M+k-j} \approx R_{yx}(i-j)$$

$$\frac{1}{T}(\mathbf{A}_X^\mathsf{T}\mathbf{A}_Y)_{i,j} = \frac{1}{T}\sum_k\sum_k x_{M+k-i}y_{M+k-j} \approx R_{xy}(i-j)$$

$$\frac{1}{T}(\mathbf{A}_X^\mathsf{T}\mathbf{A}_X)_{i,j} = \frac{1}{T}\sum_k\sum_k x_{M+k-i}x_{M+k-j} \approx R_{xx}(i-j).$$

Also,

$$\frac{1}{T}(\mathbf{A}_Y^\mathsf{T}\mathbf{y})_i = \frac{1}{T}\sum_k\sum_k y_{M+k-i}y_{M+k} \approx R_{yy}(M-i).$$

3

5. (a) Define

$$
\mathbf{A} = \begin{bmatrix}
\cos(\Omega_1(0)) & \cdots & \cos(\Omega_L(0)) & \sin(\Omega_1(0)) & \cdots & \sin(\Omega_L(0)) \\
\cos(\Omega_1(1)) & \cdots & \cos(\Omega_L(1)) & \sin(\Omega_1(1)) & \cdots & \sin(\Omega_L(1)) \\
\vdots & \cdots & \vdots & \vdots & \cdots & \vdots \\
\cos(\Omega_1(T-1)) & \cdots & \cos(\Omega_L(T-1)) & \sin(\Omega_1(T-1)) & \cdots & \sin(\Omega_L(T-1))
\end{bmatrix}
$$

and

$$
\mathbf{x} = \begin{bmatrix} x_0 \\ \vdots \\ x_{T-1} \end{bmatrix}, \quad
\boldsymbol{\beta} = \begin{bmatrix} a_1 \\ \vdots \\ a_L \\ b_1 \\ \vdots \\ b_L \end{bmatrix}.
$$

Then $\mathbf{x} \approx \mathbf{A}\boldsymbol{\beta}$.

(b) If the frequencies Ω_ℓ are not known, then the model is nonlinear.

6. *Python broadcasting.*

(a) This is best done by using slices on each of the columns of X.

```
yhat = beta[0]*X[:,0] + beta[1]*X[:,1] + beta[2]*X[:,1]*X[:,2]
```

(b) There are two ways to do this. One method is:

```
A = np.exp(−x[:,None]*beta[None,:])
yhat = np.sum(A*alpha[None,:], axis=1)
```

In this method, we first create a matrix `A[i,j]` = `np.exp(−beta[j]*x[i])`. Then, the second line multiplies each column `j` by `alpha[j]` and sums over the columns. We could also do the summation via a matrix vector product:

```
A = np.exp(−x[:,None]*beta[None,:])
yhat = A.dot(alpha)
```

(c) We can do this with the code:

```
Dxy = x[:,None,:] − y[None,:,:]
dist = np.sum(Dxy**2,axis=2)
```

To understand this, note that

```
Dxy[i,j,k] = x[i,k] − y[j,k]
```

So, when we sum in the line `np.sum(Dxy**2,axis=2`, we are summing over the index `k`.