

# Introduction to Machine Learning

## Problems: LASSO and Model Selection

Prof. Sundeep Rangan

1. *Exhaustive search.* In this problem, we will look at how to exhaustively search over all possible subsets of features. You are given three python functions:

```
model = LinearRegression() # Create a linear regression model object
model.fit(X,y)              # Fits the model
yhat = model.predict(X)     # Predicts targets given features
```

Given training data  $\mathbf{X}_{tr}, \mathbf{y}_{tr}$  and test data  $\mathbf{X}_{ts}, \mathbf{y}_{ts}$ , write a few lines of python code to:

- (a) Find the best model using only one feature of the data (i.e. one column of  $\mathbf{X}_{tr}$  and  $\mathbf{X}_{ts}$ ).
  - (b) Find the best model using only two features of the data (i.e. two columns of  $\mathbf{X}_{tr}$  and  $\mathbf{X}_{ts}$ ).
  - (c) Suppose we wish to find the best  $k$  of  $p$  features via exhaustive searching over all possible subsets of features. How many times would you need to call the `fit` function? What if  $k = 10$  and  $p = 1000$ ?
2. *Selecting a regularizer.* Suppose we fit a regularized least squares objective,

$$J(\mathbf{w}) = \sum_{i=1}^N (y_i - \hat{y}_i)^2 + \lambda \phi(\mathbf{w}),$$

where  $\hat{y}_i$  is some prediction of  $y_i$  given the model parameters  $\mathbf{w}$ . For each case below, suggest a possible regularization function  $\phi(\mathbf{w})$ . There is no single correct answer.

- (a) All parameters vectors  $\mathbf{w}$  should be considered.
- (b) Negative values of  $w_j$  are unlikely (but still possible).
- (c) For each  $j$ ,  $w_j$  should not change that significantly from  $w_{j-1}$ .
- (d) For most  $j$ ,  $w_j = w_{j-1}$ . However, it can happen that  $w_j$  can be different from  $w_{j-1}$  for a few indices  $j$ .

Variable	Units	Mean	Std dev
Median income, $x_1$	\$	50000	15000
Median age, $x_2$	years	45	10
House sale price, $y$	\$1000	300	100

Table 1: Features for Problem 3

3. *Normalization.* A data analyst for a real estate firm wants to predict house prices based on two features in each zip code. The features are shown in Table 1. The agent decides to use a linear model,

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2, \quad z_i = \frac{x_i - \bar{x}_i}{\sigma_j}. \quad (1)$$

- (a) What is the problem in using a LASSO regularizer of the form,

$$\phi(\boldsymbol{\beta}) = \sum_{j=1}^2 |\beta_j|.$$

- (b) To uniformly regularize the features, she fits a model on the normalized features,

$$\hat{u} = \alpha_1 z_1 + \alpha_2 z_2, \quad z_i = \frac{z_j - \bar{z}_j}{\sigma_j}, \quad u = \frac{\hat{y} - \bar{y}}{\sigma_y}$$

She obtains parameters  $\boldsymbol{\alpha} = [0.6, -0.3]$ ? What are the parameters  $\boldsymbol{\beta}$  in the original model (1).

4. *Normalization in python.* You are given python functions,

```
model = SomeModel()           # Creates a model
model.fit(Z,u)                 # Fits the model, expecting normalized features
yhat = model.predict(Z)        # Predicts targets given features
```

Given training data  $\mathbf{x}_{tr}, \mathbf{y}_{tr}$  and test data  $\mathbf{x}_{ts}, \mathbf{y}_{ts}$ , write python code to:

- Normalize the training data to remove the mean and standard deviation from both  $\mathbf{x}_{tr}$  and  $\mathbf{y}_{tr}$ .
- Fit the model on the normalized data.
- Predict the values  $\mathbf{y}_{hat}$  on the test data.
- Measure the RSS on the test data.

5. *Discretization.* Suppose we wish to fit a model,

$$y \approx \hat{y} = \sum_{j=1}^K \beta_j e^{-\alpha_j x}, \quad (2)$$

for parameters  $\alpha_j$  and  $\beta_j$ . Since the parameters  $\alpha_j$  are not known, this model is nonlinear and cannot be fit with least squares. A common approach in such circumstances is to use an alternate linear model,

$$y \approx \hat{y} = \sum_{j=1}^p \tilde{\beta}_j e^{-\tilde{\alpha}_j x}, \quad (3)$$

where the values  $\tilde{\alpha}_j$  are a fixed, large number  $p$  of possible values for  $\alpha_j$  and  $\tilde{\beta}_j$  are the coefficients in the model. The values  $\tilde{\alpha}_j$  are *fixed*, so only the parameters  $\tilde{\beta}_j$  need to be learned. Hence, the model (3) is linear. The model (3) is equivalent to (2) if only a small number  $K$  of the coefficients  $\tilde{\beta}_j$  are non-zero. You are given three python functions:

```

model = Lasso(lam=lam)           # Creates a linear LASSO model
                                  # with a regularization lamb
beta = model.fit(Z,y)            # Finds the model parameters using the
                                  # LASSO objective
                                  # ||y-Z*beta||^2 + lamb*||beta||_1
yhat = model.predict(Z)          # Predicts targets given features Z:
                                  # yhat = Z*beta

```

Note this syntax is slightly different from the `sklearn` syntax. You are also given training data `xtr,ytr` and test data `xts,yts`. Write python code to:

- Create  $p = 100$  values of  $\tilde{\alpha}_j$  uniformly in some interval  $\tilde{\alpha}_j \in [a, b]$  where  $a$  and  $b$  are given.
  - Fit the linear model (3) on the training data for some given `lam`.
  - Measure the test error.
  - Find coefficients  $\alpha_j$  and  $\beta_j$  corresponding to the largest  $k = 3$  values in  $\tilde{\beta}_j$ . You can use the function `np.argsort`.
6. *Minimizing an  $\ell_1$  objective.* In this problem, we will show how to minimize a simple scalar function with an  $\ell_1$ -term. Given  $y$  and  $\lambda > 0$ , suppose we wish to find the minimum,

$$\hat{w} = \arg \min_w J(w) = \frac{1}{2}(y - w)^2 + \lambda|w|.$$

Write  $\hat{w}$  in terms of  $y$  and  $\lambda$ . Since  $|w|$  is not differentiable everywhere, you cannot simply set  $J'(w) = 0$  and solve for  $w$ . Instead, you have to look at three cases:

- (i) First, suppose there is a minima at  $w > 0$ . In this region,  $|w| = w$ . Since the set  $w > 0$  is open, at any minima  $J'(w) = 0$ . Solve for  $w$  and test if the solution indeed satisfies  $w > 0$ .
- (ii) Similarly, suppose  $w < 0$ . Solve for  $J'(w) = 0$  and test if the solution satisfies the assumption that  $w < 0$ .
- (iii) If neither of the above cases have a minima, then the minima must be at  $w = 0$ .