

EE-UY/CS-UY 4563: Introduction to Machine Learning

Midterm 1, Fall 2017

1. (34 points total)

- (a) (8 points) We can use the model, $f(x) = \beta_1 x + \beta_2 x^2$. We do not have a constant term to insure that $f(0) = 0$.
- (b) (9 points) For $x \geq 100$, we want $f(x) = \text{constant}$. But, since $f(x)$ is continuous, we must have $f(x) = f(100)$ for $x \geq 100$. Thus, we can use the function,

$$f(x) = \begin{cases} \beta_1 x + \beta_2 x^2 & \text{if } x \leq 100, \\ \beta_1(100) + \beta_2(100)^2 & \text{if } x > 100. \end{cases} \quad (1)$$

In this way, the function is quadratic for $x \leq 100$ and it takes the value $f(100)$ for all $x \geq 100$. If you give the answer (1), you will get full credit. But, there are multiple ways you can write the function. For example, you can write the function in (1) as

$$f(x) = \beta_1 z + \beta_2 z^2, \quad z = \min(x, 100). \quad (2)$$

The only requirement to get full credit is that the function satisfies the required condition and has only two parameters.

- (c) (8 points) Using the parametrization in (2), the matrix \mathbf{A} should be

$$\mathbf{A} = \begin{bmatrix} z_1 & z_1^2 \\ z_2 & z_2^2 \\ z_3 & z_3^2 \\ \vdots & \vdots \end{bmatrix} = \begin{bmatrix} 20 & 20^2 \\ 70 & 70^2 \\ 100 & (100)^2 \\ \vdots & \vdots \end{bmatrix},$$

The response vector is the difference of heart rates,

$$\mathbf{y} = \begin{bmatrix} 80 - 70 \\ 90 - 75 \\ 85 - 55 \\ \vdots \end{bmatrix} = \begin{bmatrix} 10 \\ 25 \\ 30 \\ \vdots \end{bmatrix}$$

- (d) (9 points) One simple code is:

```
ntr = 50 # number of training samples
z = np.minimum(dosage[:ntr], 100)
Atr = np.column_stack((z, z**2))
ytr = hr_before[:ntr] - hr_after[:ntr]
```




Figure 1: True function (blue) and estimated linear function (green).

2. (33 points total)

- (a) (6 points) Since the true data is exponential, it is not in the model class which is linear.
- (b) (9 points) The linear approximation is

$$y_i = f_0(x_i) \approx f_0(x_{\text{nom}}) + f'_0(x_{\text{nom}})(x_i - x_{\text{nom}}).$$

This fits the linear model with

$$\begin{aligned}\beta_1 &= f'(x_{\text{nom}}) = -CD e^{-Dx_{\text{nom}}} \\ \beta_0 &= f_0(x_{\text{nom}}) - x_{\text{nom}} f'(x_{\text{nom}}) = C(1 - Dx_{\text{nom}}) e^{-Dx_{\text{nom}}}.\end{aligned}$$

Now, when the true model is linear and there is no noise, linear estimation will always obtain the ‘true parameters. Hence, we will obtain $\hat{\beta} = (\beta_0, \beta_1)$ as above.

- (c) (9 points) See Fig. 1. Note that the true and estimated functions match exactly at $x_{\text{nom}} = 2$. The bias grows as x moves away from x_{nom} , so the largest bias would occur at $x = 4$.
- (d) (9 points) One possible solution is as follows:

```
# Split into training and test
ntr = 50
xtr = x[:ntr]
ytr = y[:ntr]
xts = x[ntr:]
yts = y[ntr:]

# Loop over different model orders
dtest = range(1,6)
rss = []
for d in dtest:
    bethat = fit(xtr,ytr,d)          # Fit the model on the training data
    yhat = predict(xts,bethat)      # Predict on test data
    rss_i = np.sum((yts - yhat)**2) # Measure the test RSS
    rss.append(rss_i)
```




Figure 2: Scatter plot of the data points where the green circles are $y_i = 1$ (detected) and red circles are $y_i = 0$ (not detected). The dotted line is the boundary of a potential linear classifier.

```
# Find the minimum test RSS
iopt = np.argmin(rss)
dopt = dtest[iopt]
```

3. (33 points)

- (a) (8 points) See Fig. 2. The points are most easily drawn by recognizing that P is the radius and θ is the angle from the x_1 -axis. For example, the two green points are at ± 45 degrees with radius 3.
- (b) (8 points) You can see that they are easily linearly separable. For example, we can use a linear classifier on a vertical line,

$$\hat{y} = \begin{cases} 1 & \text{if } x_1 \geq a, \\ 0 & \text{if } x_1 < a, \end{cases}$$

for a suitable threshold level a . To select a , we need that $x_1 \geq a$ for the two green circles,

$$a \leq P \cos(\theta) = 3 \cos(\pm 45) = \frac{3}{\sqrt{2}}.$$

We also need that a is to the right of the red circle at $(P, \theta) = (1, 0)$ so

$$a \geq 1.$$

So, the linear classifier will separate the classes as long as $a \in (1, 3/\sqrt{2})$. If you selected any single value of a that works, or any other line that works, you will get full credit.

- (c) (8 points) The z score is given by

$$z = \beta_0 + \beta_1 P \cos(\theta) + \beta_2 P \sin(\theta).$$

To maximize $P(y = 1|\mathbf{x})$, we want to maximize z . Take the derivative with respect to θ :

$$z' = -\beta_1 P \sin(\theta) + \beta_2 P \cos(\theta) = 0 \Rightarrow \theta = \tan^{-1}(\beta_2/\beta_1).$$

You will get full credit for this solution. But, note that there is another point where $z' = 0$: $\theta = \tan^{-1}(\beta_2/\beta_1) + \pi$. You can verify that this point will give the smallest z , (i.e. the minimum probability of detection). But, you do not need to discuss this.

(d) (9 points) One solution is:

```
def predict(P,theta,beta):  
    x1 = P*np.cos(theta)  
    x2 = P*np.sin(theta)  
    z = beta[0] + beta[1]*x1 + beta[2]*x2  
    py = 1/(1+np.exp(-z))  
    return py
```