

EE-UY/CS-UY 4563: Introduction to Machine Learning

Midterm 2, Fall 2017

Answer all THREE questions. Exam is closed book. No electronic aids. But, you are permitted a limited number of cheat sheets. Part marks are given. If you do not remember a particular python command or its syntax, use pseudo-code and state what syntax you are assuming.

Best of luck!

1. Given a time-sequence $x_i, i = 0, 1, \dots, T - 1$, suppose we try to fit a model,

$$x_i \approx \hat{x}_i = \phi \left(\sum_{j=1}^p w_j x_{i-j} \right), \quad \phi(z_i) = \frac{1}{1 + e^{-z_i}},$$

where $\mathbf{w} = (w_1, \dots, w_p)$ are the unknown parameters. Assume $x_k = 0$ for $k < 0$. To fit the model, we use a loss function,

$$J(\mathbf{w}) = \sum_{i=1}^{T-1} (\hat{x}_i - x_i)^2.$$

- (a) Find a matrix \mathbf{A} and vector \mathbf{z} such that $\mathbf{z} = \mathbf{A}\mathbf{w}$ and

$$\hat{x}_i = \phi(z_i), \quad i = 1, \dots, T - 1.$$

- (b) What is the gradient $\nabla_{\mathbf{w}} J(\mathbf{w})$?

- (c) Suppose that $\mathbf{w}^0 = (w_1^0, \dots, w_p^0)$ and $\mathbf{w}^1 = (w_1^1, \dots, w_p^1)$ are two parameter vectors such that,

$$\begin{aligned} w_1^1 &= w_1^0 + \delta, \\ w_j^1 &= w_j^0, \text{ for } j = 2, 3, \dots, p. \end{aligned}$$

That is, the two parameters are equal except for the first coordinate. Approximately, what is $J(\mathbf{w}^1) - J(\mathbf{w}^0)$ when δ is small? Use a linear approximation and leave your answer in terms of the coefficients of the gradient $\nabla J(\mathbf{w})$.

- (d) Write a short python function to implement the following variant of gradient descent:

$$\begin{aligned} \mathbf{v}^k &= \mathbf{w}^k - \alpha \nabla_{\mathbf{w}} J(\mathbf{w}^k), \\ \mathbf{w}^{k+1} &= \max\{0, \mathbf{v}^k\}. \end{aligned}$$

In the second step, the maximum is applied elementwise, $\hat{w}_j^{k+1} = \max\{0, v_j^k\}$. This algorithm is called *projected gradient descent*.

Your function should work for an *arbitrary* loss function $J(\mathbf{w})$ – not necessarily one of the loss functions above. To do this, use the format,

```
def proj_grad(feval, ...):
    ...
```

which takes a function argument `feval` that returns the loss function and gradient (i.e. `J, Jgrad = feval(w)`). Add any other arguments and make any other assumptions as necessary. The function should return the final estimate \mathbf{w}^k and loss function $J(\mathbf{w}^k)$.

2. You are given four data points \mathbf{x}_i with binary class labels $y_i = \pm 1$:

x_{i1}	0	0	2	3
x_{i2}	0	2	2	0
y_i	-1	1	-1	1

- (a) Draw a scatter plot of the four data points indicating the two classes in different markers.
(b) Find a weight, $\mathbf{w} = (w_1, w_2)$, and bias, b , such that the linear classifier,

$$\hat{y} = \begin{cases} 1, & \text{if } z > 0 \\ -1 & \text{if } z < 0. \end{cases} \quad z = b + w_1x_1 + w_2x_2.$$

makes a minimum number of errors on the training data.

- (c) Consider the SVM loss,

$$J(\mathbf{w}, b) := C \sum_{i=1}^N \epsilon_i + \frac{1}{2} \|\mathbf{w}\|^2, \quad \epsilon_i = \max\{0, 1 - y_i z_i\},$$

for some $C > 0$. For your classifier, which sample (\mathbf{x}_i, y_i) has the largest ϵ_i ? What is the value of ϵ_i ?

Note: You do not need to compute ϵ_i for all the samples. Think about the sample that will have the highest value.

- (d) Now, consider an SVM classifier,

$$\hat{y} = \begin{cases} 1, & \text{if } z > 0 \\ -1 & \text{if } z < 0, \end{cases} \quad z = \sum_{i=1}^N \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}),$$

using a radial basis function,

$$K(\mathbf{x}_i, \mathbf{x}) = e^{-\gamma \|\mathbf{x}_i - \mathbf{x}\|^2}, \quad \|\mathbf{x}_i - \mathbf{x}\|^2 = \sum_{j=1}^d (x_{ij} - x_j)^2.$$

Write a python method `predict`, that outputs \hat{y} for a single data point \mathbf{x} :

```
def predict(x,...):
    return yhat
```

You will need to supply your method `predict` any arguments in addition to \mathbf{x} that you will need. You can assume the input dimension is $d = 2$.

3. Consider a neural network used for binary classification of the form,

$$z_j^H = \sum_{k=1}^{N_i} W_{jk}^H x_k + b_j^H, \quad u_j^H = \begin{cases} 1, & \text{if } z_j^H > 0, \\ 0, & \text{if } z_j^H < 0. \end{cases}, \quad j = 1, \dots, N_h$$

$$z^O = \sum_{k=1}^{N_h} W_k^O u_k^H + b^O, \quad P(y = 1|\mathbf{x}) = \frac{1}{1 + e^{-z^O}}.$$

The hidden weights and biases are:

$$\mathbf{W}^H = \begin{bmatrix} 1 & 0 \\ -1 & 0 \\ 0 & 1 \end{bmatrix}, \quad \mathbf{b}^H = \begin{bmatrix} -2 \\ 4 \\ -1 \end{bmatrix}.$$

- Draw the region on the (x_1, x_2) plane where $u_j^H = 1$ for all j .
- Suppose that the output weight vector is $\mathbf{W}^O = [1, 1, 1]$. For what range of values b^O is $\hat{y} = 1$ when $\mathbf{x} = (3, 2)$.
- Consider the problem of computing the gradient of a loss function J on a mini-batch (\mathbf{x}_i, y_i) , $i = 1, \dots, N$. In back-propagation, suppose that we have computed the gradients, $\partial J / \partial z_{ij}^H$. Show how to compute the gradients $\partial J / \partial W_{jk}^H$.
- Write a few lines of python code to implement the gradient calculation in part (c). State your assumptions on how you represent the inputs and outputs for the calculation.