

EE-UY/CS-UY 4563: Introduction to Machine Learning
Midterm 2 Solutions, Fall 2017

1. (33 points)

(a) (8 points) Let

$$z_i = \sum_{j=1}^p x_{i-j} w_j,$$

so that $\hat{x}_i = \phi(z_i)$. Also, $\mathbf{z} = \mathbf{A}\mathbf{w}$ if

$$\mathbf{A} = \begin{bmatrix} x_0 & 0 & \cdots & 0 \\ x_1 & x_0 & \cdots & 0 \\ x_{T-2} & x_{T-2} & \cdots & x_{T-p-1} \end{bmatrix}.$$

(b) (8 points) The loss function is

$$J(\mathbf{w}) = \sum_{i=1}^{T-1} g_i(z_i), \quad g_i(z_i) = (\phi(z_i) - x_i)^2.$$

Using the forward-backward rule,

$$\nabla J(\mathbf{w}) = \mathbf{A}^\top \nabla_{\mathbf{z}} g(\mathbf{z}),$$

where

$$\nabla_{\mathbf{z}} g(\mathbf{z}) = [g'_1(z_1), \dots, g'_{T-1}(z_{T-1})]^\top,$$

and

$$g'_i(z_i) = 2(\phi(z_i) - x_i)\phi'(z_i) = 2(\phi(z_i) - x_i) \frac{e^{-z_i}}{1 - e^{-z_i}}.$$

(c) (9 points) The first order approximation is

$$\begin{aligned} J(\mathbf{w}^1) - J(\mathbf{w}^0) &\approx \nabla J(\mathbf{w}^0) \cdot (\mathbf{w}^1 - \mathbf{w}^0) \\ &= \sum_{j=1}^p \frac{\partial J(\mathbf{w}^0)}{\partial w_j} (w_j^1 - w_j^0) = \delta \frac{\partial J(\mathbf{w}^0)}{\partial w_1}. \end{aligned}$$

(d) (9 points) One possible implementation is given as:




Figure 1: Scatter plot of the data points where the green circles are $y_i = 1$ and the red circles are $y_i = -1$. The dotted line is the boundary of a potential linear classifier.

```
def proj_grad(feval, alpha, nit, winit):
    w = winit
    for it in range(nit):
        J, Jgrad = feval(w)
        v = w - alpha * Jgrad
        w = np.maximum(0, v)
    return w, J
```

2. (33 points total)

(a) (8 points) See Fig. 1.

(b) (8 points) The data is not linearly separable. But, if we take the classifier with $z = x_1 + x_2 - 1$, this makes only one error. The weights and bias for this classifier is:

$$\mathbf{w} = [1, 1], \quad b = -1.$$

Of course, there are other classifiers that make only one error. Any such choice will receive full marks, but you must specify \mathbf{w} and b and not simply draw the line.

(c) (8 points) The point at $(2, 2)$ is misclassified by the classifier in part (b). Since this is the only misclassified point, it will have the highest hinge loss. For this point,

$$z_i = w_1 x_{i1} + w_2 x_{i2} + b = (1)(2) + (1)(2) - 1 = 3.$$

Hence, the hinge loss is

$$\epsilon_i = \max\{0, 1 - y_i z_i\} = \max\{0, 1 - (-1)(3)\} = 4.$$

(d) (9 points) A simple implementation is as follows:

```
def predict(x, xtr, ytr, alpha, gam):
    dsq = (x[0] - xtr[:, 0])**2 + (x[1] - xtr[:, 1])**2
    z = np.sum(ytr * alpha * np.exp(-gam * dsq))
    yhat = (z > 0)
```




Figure 2: Region where $u_j^H = 1$ for all j is the described by the intersection of the three half-planes.

Note that the function requires the training data `xtr,ytr` as well as `alpha` and `gam`.

You will receive full credit for the above code. But, if you wanted the code to work for an arbitrary dimension d , you could compute the distance `dsq` with the following code which uses python broadcasting.

```
dsq = np.sum((x[None,:] - xtr)**2, axis=1)
```

The other lines would remain the same.

3. (34 points total)

(a) (9 points) Since $\mathbf{z}^H = \mathbf{W}^H \mathbf{x} + \mathbf{b}^H$,

$$\mathbf{z}^H = \begin{bmatrix} z_1^H \\ z_2^H \\ z_3^H \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ -1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + \begin{bmatrix} -2 \\ 4 \\ -1 \end{bmatrix} = \begin{bmatrix} x_1 - 2 \\ -x_1 + 4 \\ x_2 - 1 \end{bmatrix}.$$

For each hidden output j , $u_j^H = 1$ when $z_j^H > 0$. Therefore,

$$\begin{aligned} u_1^H = 1 &\iff x_1 - 2 > 0 \iff x_1 > 2 \\ u_2^H = 1 &\iff -x_1 + 4 > 0 \iff x_1 < 4 \\ u_3^H = 1 &\iff x_2 - 1 > 0 \iff x_2 > 1. \end{aligned}$$

Thus, each region where $u_j^H = 1$ is a half-plane and their intersection is shown in Fig. 2.

(b) (9 points) Since $\mathbf{x} = (3, 2)$ is in the region in Fig. 2 we have $u_j^H = 1$ for all j , Therefore with $\mathbf{W}^O = (1, 1, 1)$,

$$z^O = \sum_{j=1}^3 W_j^O u_j^H + b^O = (1)(1) + (1)(1) + (1)(1) + b^O = 3 + b^O.$$

In order that $\hat{y} = 1$, we need $z^O > 0$ and therefore,

$$z^O > 0 \iff 3 + b^O > 0 \iff b^O > -3.$$

(c) (8 points) We have

$$z_{ij}^H = \sum_{k=1}^{N_h} W_{jk}^H x_{ik} + b_j^H \implies \frac{\partial z_{ij}^H}{\partial W_{jk}^H} = x_{ik}$$

By chain rule,

$$\frac{\partial J}{\partial W_{jk}^H} = \sum_{i=1}^N \frac{\partial J}{\partial z_{ij}^H} \frac{\partial z_{ij}^H}{\partial W_{jk}^H} = \sum_{i=1}^N \frac{\partial J}{\partial z_{ij}^H} x_{ik}.$$

(d) (8 points) We assume that x_{ij} is represented in a matrix \mathbf{x} and $\partial J / \partial z_{ij}^H$ is represented in a matrix $\mathbf{Jgrad_zhid}$. Then, we can perform the sum in part (c) as:

```
Jgrad_Whid = np.zeros((nhid,nin))
for j in range(nhid):
    for k in range(nin):
        Jgrad_Whid[j,k] = np.sum(zhid[:,j]*x[:,k])
```

This will get full credit. But, you can avoid the for loops using python broadcasting:

```
Jgrad_Whid = np.sum(zhid[:, :, None] * x[:, None, :], axis=0)
```